



GENERATIVE ADVERSARIAL NETWORKS

Deep Neural Networks

Session 23

Pramod Sharma


pramod.sharma@prasami.com

2

Agenda



PixelRNN / CNN



Variational Autoencoder



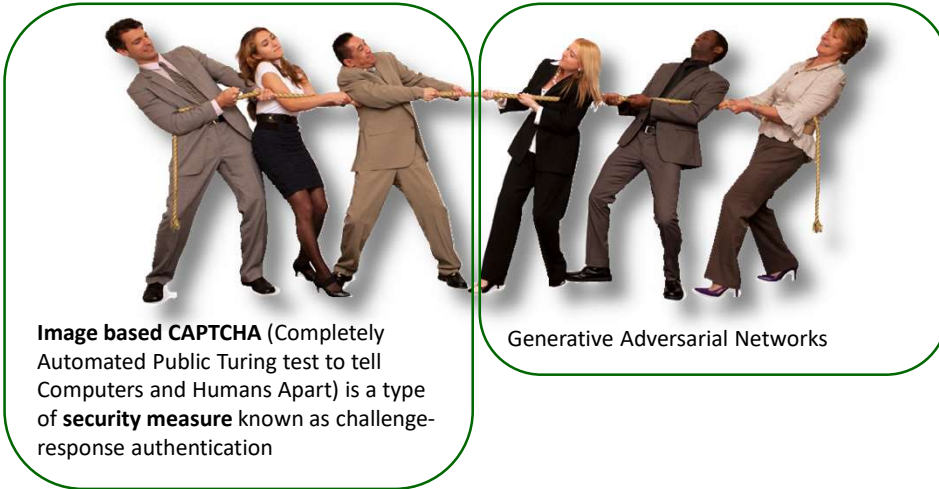
GAN

6/1/2024

pra-sâmi

3

Technologies with Conflicting Goals




6/1/2024

pra-sâmi

4

Unbelievably Real

- ❑ This person does not exist : thispersondoesnotexist.com



6/1/2024

pra-sâmi

5

Overview

- ❑ Developed by Ian Goodfellow
- ❑ In generative modeling, we'd like to train a network that models a distribution,
 - ❖ Such as a distribution over images.
- ❑ One way to judge the quality of the model is to sample from it
- ❑ Active area of research with rapid progress



2009



2015



2018

6/1/2024

pra-sâmi

6

Take a Step Back

Supervised Learning

- ❑ Data: (x, y)
 - ❖ x is data, y is label
- ❑ Goal: Learn a function to map $x \rightarrow y$
- ❑ Examples:
 - ❖ Classification,
 - ❖ Regression,
 - ❖ Object detection,
 - ❖ Semantic segmentation,
 - ❖ Image captioning,
 - ❖ ...

Unsupervised Learning

- ❑ Data: x
 - ❖ Just data, no labels!
- ❑ Goal: Learn some underlying hidden structure of the data
- ❑ Examples:
 - ❖ Clustering,
 - ❖ Dimensionality reduction,
 - ❖ Feature learning,
 - ❖ Density estimation,
 - ❖ ...

Makes Training
data cheap!

Holy grail:
Solve unsupervised learning
→ Understand structure of
visual world

6/1/2024

pra-sâmi

7

Generative Models

- Given the training data, generate new samples from same distribution



Training data $\sim p_{\text{data}}(x)$



Generated samples $\sim p_{\text{model}}(x)$

- Want to learn $p_{\text{model}}(x)$ similar to $p_{\text{data}}(x)$
- Addresses density estimation, a core problem in unsupervised learning
- Several flavors:
 - ❖ Explicit density estimation: explicitly define and solve for $p_{\text{model}}(x)$
 - ❖ Implicit density estimation: learn model that can sample from $p_{\text{model}}(x)$ without explicitly defining it

6/1/2024

pra-sâmi

8

Why Generative Models?

- Realistic samples for artwork, super-resolution, colorization, etc.



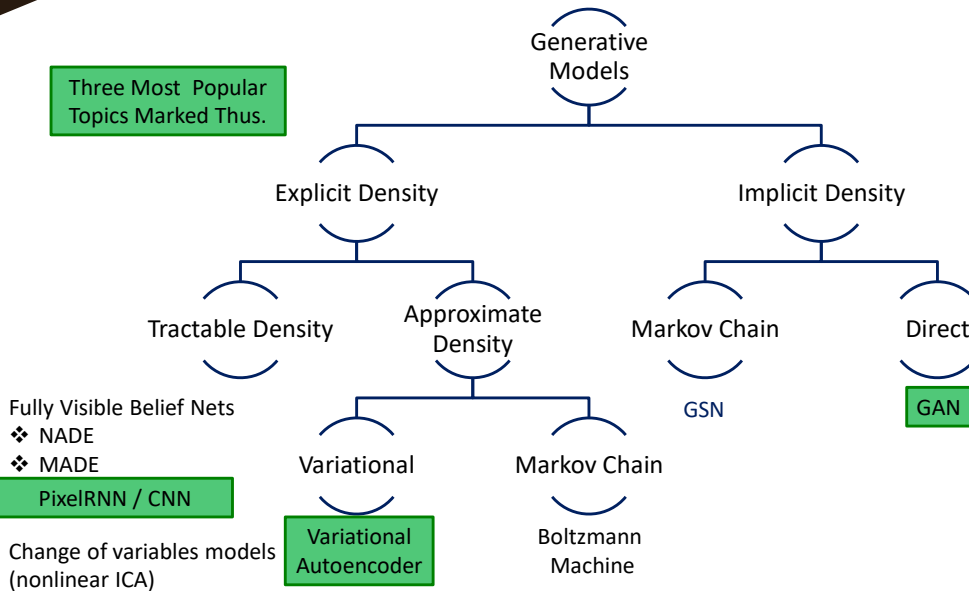
- Generative models of time-series data can be used for simulation and planning
 - ❖ Such as reinforcement learning applications!
- Training generative models can also enable inference of latent representations that can be useful as general features

6/1/2024

pra-sâmi

9

Taxonomy of Generative Models



10

Fully Visible Belief Network

- ❑ Explicit Density Model
- ❑ Use chain rule to decompose likelihood of an image x into product of 1-d distributions:

$$p(x) = \prod_{i=1}^n p(x_i | x_1, x_2, x_3 \dots x_{i-1})$$

Likelihood
of image x

Probability of i th pixel value
given all previous pixels

- ❑ Then maximize likelihood of training data
- ❑ Note:
 - ❖ Will need to define ordering of “previous pixels”
 - ❖ Complex distribution over pixel values → Express using a neural network!


6/1/2024

pra-sâmi

11

PixelRNN [van der Oord et al. 2016]

- ❑ Generate image pixels starting from corner
- ❑ Dependency on previous pixels modeled using an RNN (LSTM)




6/1/2024

pra-sâmi

12

PixelRNN [van der Oord et al. 2016]

- ❑ Generate image pixels starting from corner
- ❑ Dependency on previous pixels modeled using an RNN (LSTM)




6/1/2024

pra-sâmi

13

PixelRNN [van der Oord et al. 2016]

- Generate image pixels starting from corner
- Dependency on previous pixels modeled using an RNN (LSTM)




6/1/2024

pra-sâmi

14

PixelRNN [van der Oord et al. 2016]

- Generate image pixels starting from corner
- Dependency on previous pixels modeled using an RNN (LSTM)




6/1/2024

pra-sâmi

15

PixelRNN [van der Oord et al. 2016]

- Generate image pixels starting from corner
- Dependency on previous pixels modeled using an RNN (LSTM)



- Drawback:
 - ❖ Very sequential generation, very slow!

6/1/2024

pra-sâmi


16

PixelCNN [van der Oord et al. 2016]

- PixelCNN also generates image pixels starting from corner,
- Dependency on previous pixels now modeled using a CNN over context region
- Training: maximize likelihood of training images

$$p(x) = \prod_{i=1}^n p(x_i | x_1, x_2, x_3 \dots x_{i-1})$$

- Training is faster than PixelRNN
 - ❖ can parallelize convolutions since context region values known from training images
- Generation must still proceed sequentially
 - ➔ still slow!



6/1/2024

pra-sâmi

17

Generation Samples



32x32 CIFAR-10



32x32 ImageNet

6/1/2024

pra-sâmi

18

PixelRNN and PixelCNN

Pros:

- ❑ Can explicitly compute likelihood $p(x)$
- ❑ Explicit likelihood of training data gives good evaluation metric
- ❑ Good samples

Con:

- ❑ Sequential generation → slow

Reference

- Van der Oord et al. NIPS 2016
- Salimans et al. 2017 (PixelCNN++)

Improving PixelCNN performance

- ❑ Gated convolutional layers
- ❑ Short-cut connections
- ❑ Discretized logistic loss
- ❑ Multi-scale
- ❑ Training tricks
- ❑ Etc...

6/1/2024

pra-sâmi

19

Overview

- ❑ Four modern approaches to generative modeling:
 - ❖ **Generative adversarial networks**
 - ❖ Reversible architectures
 - ❖ Autoregressive models
 - ❖ **Variational autoencoders**
- ❑ All four approaches have different pros and cons
- ❑ In this session we will focus on
 - ❖ Variational autoencoders i.e. VAEs
 - ❖ Generative Adversarial Networks i.e. GANs

6/1/2024

pra-sâmi

20

Variational Autoencoders (VAE)

6/1/2024

pra-sâmi

21

Difference between PixelCNN and VAE

- PixelCNNs define **tractable** density function, optimize likelihood of training data:

$$p(x) = \prod_{i=1}^n p(x_i | x_1, x_2, x_3 \dots x_{i-1})$$

- VAEs define **intractable** density function with latent z :

$$p_{\theta}(x) = \int p_{\theta}(z) p_{\theta}(x|z) dz$$

- Cannot optimize directly,
 - ❖ So we derive and optimize lower bound on likelihood instead
- Too lengthy, remained theoretical discussions...
- What if we give up on explicitly modeling density, and just want ability to sample?


6/1/2024

pra-sâmi

22

Background: Autoencoders

- Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data
- How to learn these features
 - ❖ Train such that features can be used to reconstruct original data
 - ❖ “Autoencoding” – encoding itself.



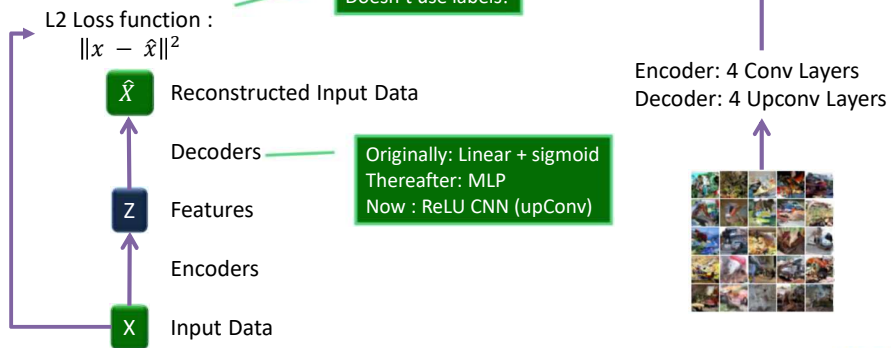
6/1/2024

pra-sâmi

23

Background: Autoencoders

- ❑ Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data
- ❑ How to learn these features
 - ❖ Train such that features can be used to reconstruct original data
 - ❖ “Autoencoding” – encoding itself.



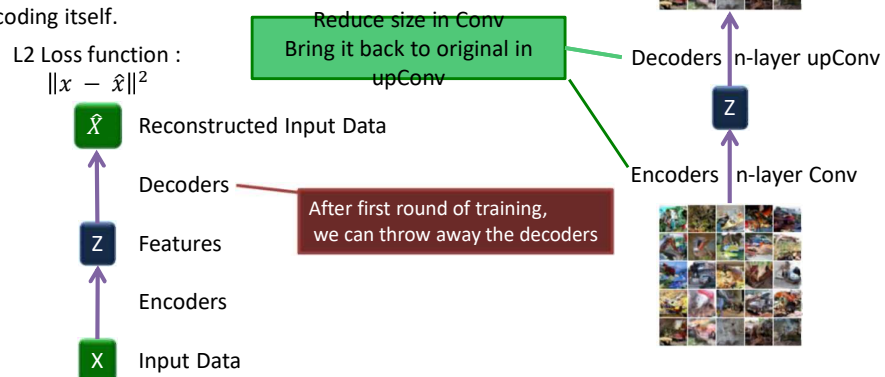
6/1/2024

pra-sâmi

24

Background: Autoencoders

- ❑ Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data
- ❑ How to learn these features
 - ❖ Train such that features can be used to reconstruct original data
 - ❖ “Autoencoding” – encoding itself.



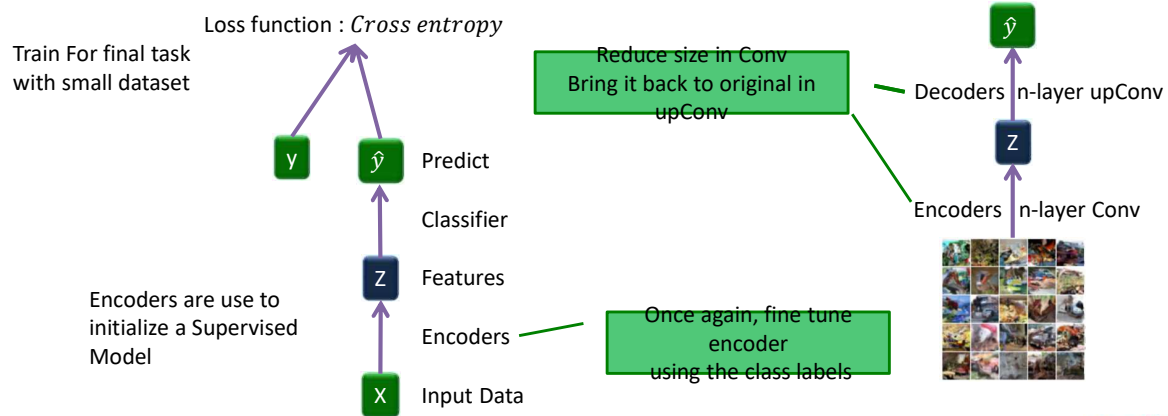
6/1/2024

pra-sâmi

25

Autoencoders

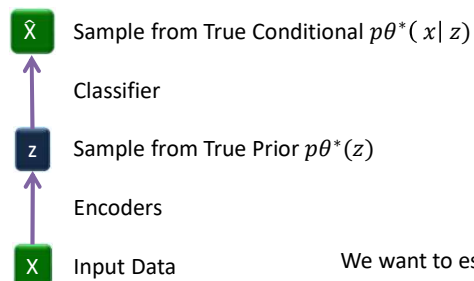
- Now job of decoder is done and we have optimized weights of the Encoder
- Use this encoder for your analysis



26

Variational Autoencoders

- Probabilistic spin on auto encoders
 - ❖ Will let us sample from the generated data
- Assume that Training data is generated from some underlying unobserved (latent) representation z



Intuition

- X is an image, z is latent factors use to generate x
- Attributes: pose, smile, pupil dilation, angle of eyebrow etc.

We want to estimate the true parameters θ^* of this generative model!

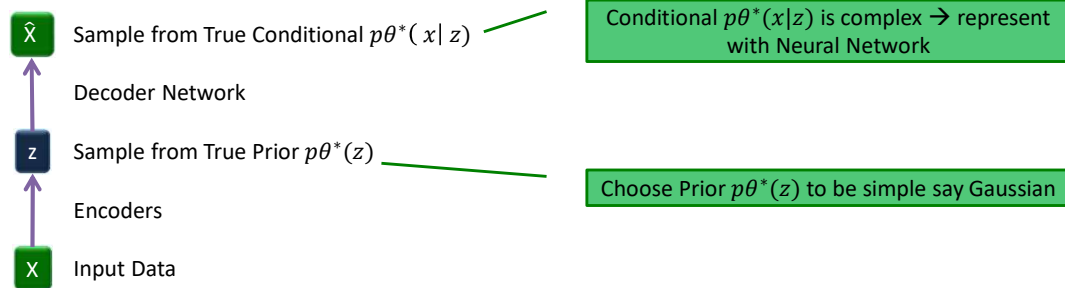
6/1/2024

pra-sâmi

27

Variational Autoencoders

- We want to estimate the true parameters θ^* of this generative model!
- How should we represent this model?



6/1/2024

pra-sâmi

28

Variational Autoencoders

- Straightforward way is to maximize likely hood of data model
 - ❖ $p_{\theta}(x) = \int p_{\theta}(z) * p_{\theta}(x|z) * dz$
 - ❖ We need to integrate as we are looking at all possible values of x
 - ❖ Hence it is not tractable.
- In details: data likelihood $p_{\theta}(x) = \int p_{\theta}(z) * p_{\theta}(x|z) * dz$
 - ❖ $p_{\theta}(z) \rightarrow$ ok, we can use Gaussian Prior probabilities
 - ❖ $p_{\theta}(x|z) \rightarrow$ Ok too as we can use a decode Neural Network
 - ❖ Integration is a problem, as we need to look at all possible values of z
- It turns out that posterior $p_{\theta}(x|z)$ is also intractable (difficult to converge)
 - ❖ $p_{\theta}(z|x) = p_{\theta}(x|z) * \frac{p_{\theta}(z)}{p_{\theta}(x)}$ Intractable
- Solution:
 - ❖ Decoder model for $p_{\theta}(x|z)$ and a separate encoder model $q_{\theta}(z|x)$

6/1/2024

pra-sâmi

31

Explicit Density Models

- ❑ PixelCNNs define tractable density function, optimize likelihood of training data:

$$p(x) = \prod_{i=1}^n p(x_i | x_1, x_2, x_3 \dots x_{i-1})$$

- ❑ VAEs define intractable density function with latent z :

$$p_{\theta}(x) = \int p_{\theta}(z) p_{\theta}(x|z) dz$$

- ❑ Cannot optimize directly, derive and optimize lower bound on likelihood instead
- ❑ Too lengthy, remained theoretical discussions...
- ❑ What if we give up on explicitly modeling density, and just want ability to sample?
- ❑ GANs: don't work with any explicit density function! Instead, take game-theoretic approach: learn to generate from training distribution through 2-player game

6/1/2024

pra-sâmi

32

Generative Adversarial Networks


6/1/2024

pra-sâmi


33

Generative Adversarial Networks

- The idea behind Generative Adversarial Networks (GANs): train two different networks
 - ❖ The generator network tries to produce realistic-looking samples
 - ❖ The discriminator network tries to figure out whether an image came from the training set or the generator network
- The generator network tries to fool the discriminator network



Generator




Discriminator

6/1/2024


pra-sâmi

34

Generative Adversarial Network



Generator



Discriminator





6/1/2024

pra-sâmi

35

Generative Adversarial Network





6/1/2024

pra-sâmi

36

Generative Adversarial Network




6/1/2024

pra-sâmi

37

Generative Adversarial Network




6/1/2024

pra-sâmi

38

Generative Adversarial Network




6/1/2024

pra-sâmi

39

Generative Adversarial Network




6/1/2024

pra-sâmi

40

Generative Adversarial Network




6/1/2024

pra-sâmi

41

Generative Adversarial Network

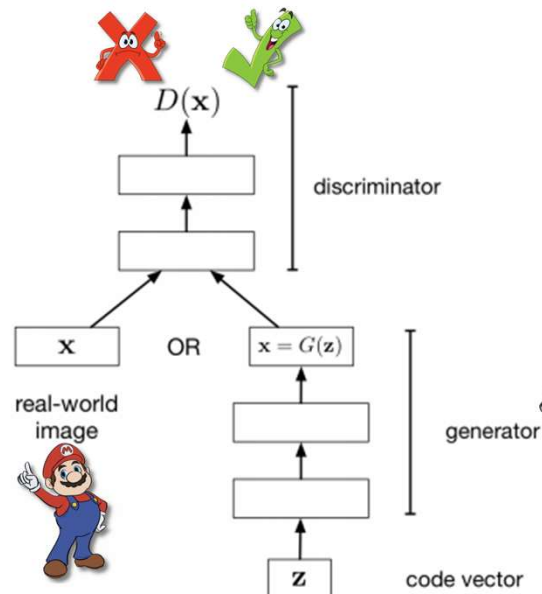


6/1/2024

pra-sâmi

42

Generative Adversarial Networks



- ❑ Discriminator network is trained on real images as well as generated images
- ❑ Generator network: try to fool the discriminator by generating real-looking images
- ❑ Discriminator network: try to distinguish between real and fake images

6/1/2024


pra-sâmi

43

Coding Gan

- ❑ Imagine a simplest 2 x 2 images
 - ❖ Depending upon features these shades may vary

- ❑ Let's also consider our simplest neural network



6/1/2024

pra-sâmi

44

GAN

- ❑ Imagine all images are slanted backward by 45°
- ❑ We have following images of faces



- ❑ The corresponding pixel on the images will be as follows:



- ❑ For argument sake let's take white pixel as 0 and black as 1
 - ❖ Gray shades will be somewhere in between



6/1/2024

pra-sâmi

45

GAN

- We have following images of faces



- Images containing no face will appear as follows:






6/1/2024

pra-sâmi





46

What Agent will see

- Faces



- No Faces



6/1/2024

pra-sâmi

47

What Agent will see

□ Faces

0.75	0
0	0.75



1	0.5
0.5	1

1	0
0	1

1	0.25
0	1

□ No Faces

0.25	0
1	0.75

0.75	1
1	0.25

0.75	0.25
0.75	0.25

1	1
0	0.25

6/1/2024

pra-sâmi

48

Discriminator

6/1/2024


pra-sâmi

49


Discriminator

How to differentiate between Face and Noise!

□ Faces



□ Noise



6/1/2024


pra-sâmi

50

Discriminator


How to differentiate between Face and Noise!

□ Faces



$$\square 1 * 1 + 1 * 1 + 0 * (-1) + 0 * (-1) = 2.0$$

□ Noise



$$\square 0.75 * 1 + 0.25 * 1 + 1 * (-1) + 1 * (-1) = -1.0$$

Considering threshold as 0.0 (arbitrary), we can safely assume:

- ❖ Positive value → Face
- ❖ Negative value → No Face


6/1/2024

pra-sâmi

51

Discriminator

In Neuron Network Language

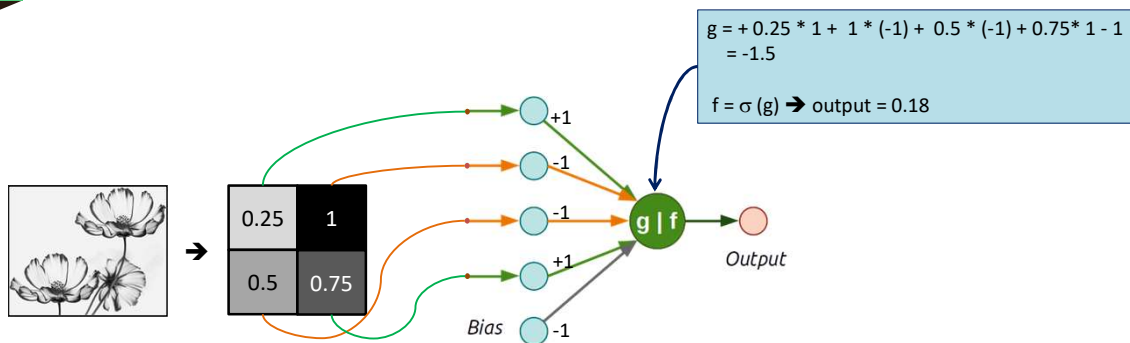


6/1/2024

pra-sâmi

52

Discriminator



□ Thus

- ❖ for $f > 0.5$ it is face
- ❖ for $f < 0.5$ it is not a face


6/1/2024

pra-sâmi

53

Discriminator

- All images are slanted backward by 45°



0.5	0
0	0.5



1	0.5
0.5	1



1	0
0	1



1	0.5
0	1



0.25	1
0.5	0.75

- So the discriminator knows that following:

❖ Face



Not a Face



6/1/2024

pra-sâmi

54


Generator

6/1/2024

pra-sâmi

55

Generator

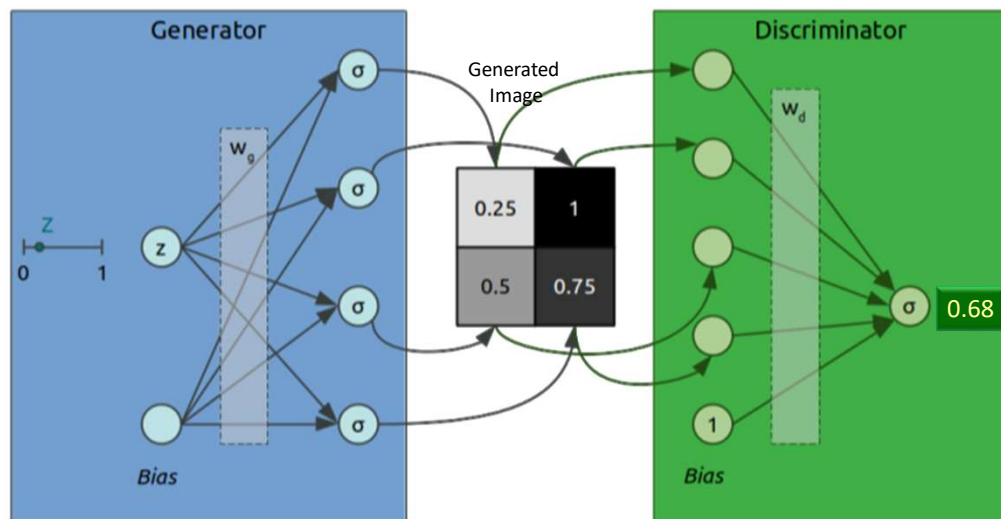


6/1/2024

pra-sâmi

56

Error Functions




6/1/2024

Generator is Max and Discriminator is Min

pra-sâmi

57

Error Functions



- ❑ Generator and Discriminators are working against each other
- ❑ Discriminator tries to generate label as close to 0 as possible (Claiming it is fake)
 - ❖ Error function = $-\log(1-p)$
- ❑ Generator tries to generate labels as close to 1 as possible (Claiming it to be an image)
 - ❖ Error function = $-\log(p)$

6/1/2024

pra-sâmi

58

Error Functions

Discriminator

- ❑ If our value 0 and prediction is 0.1 → error is small
- ❑ If our value is 0 and prediction is 0.9 → error is large
- ❑ Consider negative log error
 - ❖ For pred = 0.1; error = $-\ln(1 - 0.1) = 0.11$
 - ❖ For pred = 0.9 error = $-\ln(1 - 0.9) = 2.30$
- ❑ Thus our error function is:
 - ❖ $-\ln(1 - \text{pred})$

Generator

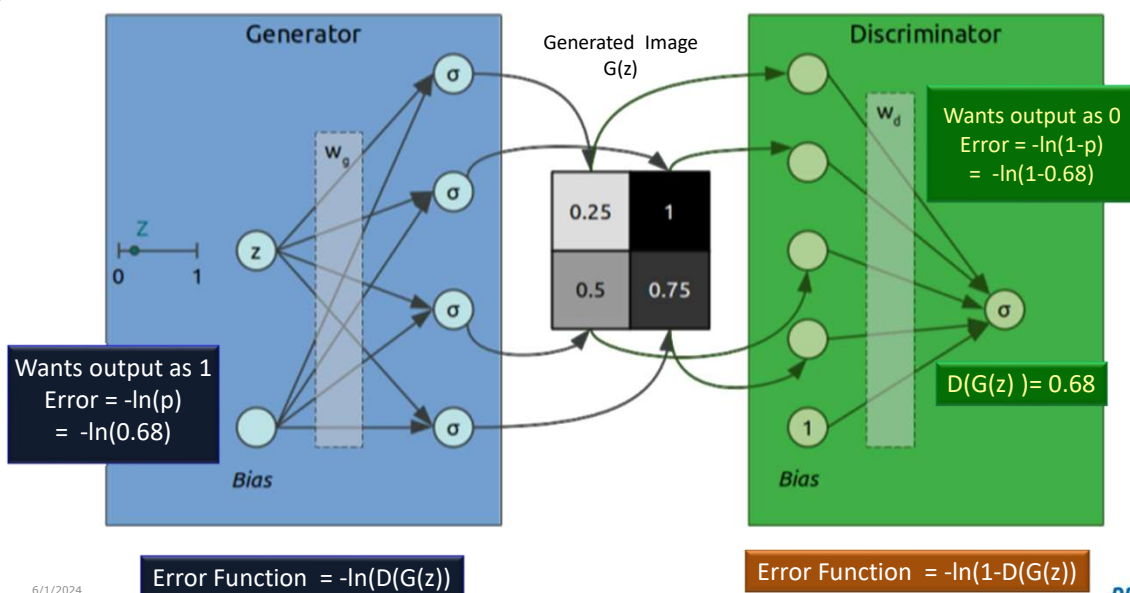
- ❑ If our value 1 and prediction is 0.1 → error is large
- ❑ If our value is 1 and prediction is 0.9 → error is small
- ❑ Consider negative log error
 - ❖ For pred = 0.1; error = $-\ln(0.1) = 2.30$
 - ❖ For pred = 0.9 error = $-\ln(0.9) = 0.1$
- ❑ Thus our error function is:
 - ❖ $-\ln(\text{pred})$

6/1/2024

pra-sâmi

59

Error Functions



60

Three Reasons that it's a Miracle GANs Work

- ❑ G has a reinforcement learning task
 - ❖ It knows when it does good (i.e., fools D) but it is not given a supervised signal
 - ❖ Reinforcement learning is hard
 - ❖ Back prop through D provides G with a supervised signal; the better D is, the better this signal will be
- ❑ Can't describe optimum via a single loss
 - ❖ Will there be an equilibrium?
- ❑ D is seldom fooled
 - ❖ But G still learns because it gets a gradient telling it how to change in order to do better the next round.

6/1/2024

pra-sâmi

61

Training GANs: Two-player game

- ❑ Generator network: try to fool the discriminator by generating real-looking images
- ❑ Discriminator network: try to distinguish between real and fake images
- ❑ Train jointly in MiniMax game
- ❑ MiniMax objective function:

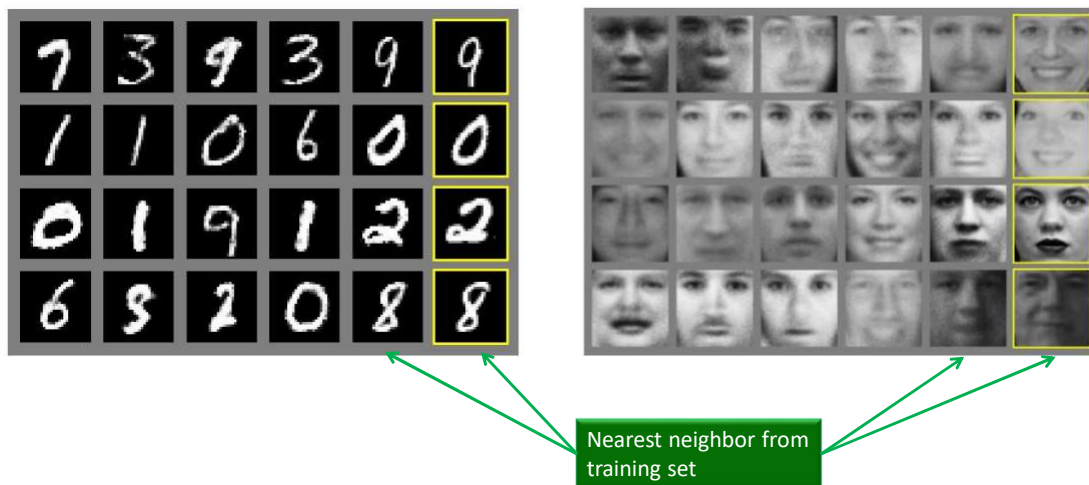
$$\min_{\theta_g} \max_{\theta_d} [E_{x \sim p_{data}} \log D_{\theta_d}(x) + E_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))]$$
- ❑ Discriminator outputs likelihood in (0,1) of real image
- ❑ Discriminator (θ_d) wants to maximize objective such that $D(x)$ is close to 1 (real) and $D(G(z))$ is close to 0 (fake)
- ❑ Generator (θ_g) wants to minimize objective such that $D(G(z))$ is close to 1
 - ❖ Discriminator is fooled into thinking generated $G(z)$ is real

6/1/2024

pra-sâmi

62

Generative Adversarial Nets



6/1/2024

pra-sâmi

64

Reflect...

- ❑ Which of the following are key components of a Generative Adversarial Network (GAN)?
 - a. Generator
 - b. Discriminator
 - c. Classifier
 - d. Loss function
- ❑ Answer : a, b, d
- ❑ Select the statements that correctly describe the training process of a GAN.
 - a. The generator aims to produce data that is indistinguishable from real data.
 - b. The discriminator provides feedback to the generator about the generated samples.
 - c. GANs are trained using supervised learning techniques.
 - d. The loss function for GANs involves both a generator loss and a discriminator loss.
- ❑ Answer: a, b, d

6/1/2024

pra-sâmi

- ❑ Which applications can benefit from the use of Generative Adversarial Networks?
 - a. Image generation
 - b. Style transfer
 - c. Text summarization
 - d. Speech recognition
- ❑ Answer : a, b, d
- ❑ What is the purpose of the generator in a GAN?
 - a. To discriminate between real and fake data.
 - b. To generate synthetic data.
 - c. To evaluate the quality of generated samples.
 - d. To provide feedback to the discriminator.
- ❑ Answer : b

65

Reflect...

- ❑ Choose the correct statements regarding the mode collapse phenomenon in GANs.
 - a. Mode collapse occurs when the generator produces diverse samples covering the entire data distribution.
 - b. Mode collapse happens when the generator focuses on generating only a limited set of samples.
 - c. Mode collapse is a desired behavior in GAN training.
 - d. Mode collapse is related to the overfitting of the discriminator.
- ❑ Answer : b
- ❑ Which regularization techniques are commonly used to stabilize GAN training?
 - a. Dropout
 - b. Batch normalization
 - c. L1 regularization
 - d. Gradient clipping
- ❑ Answer : a, b, d

6/1/2024

pra-sâmi

- ❑ Select the statements that correctly describe the challenges associated with training Generative Adversarial Networks.
 - a. GANs may suffer from mode collapse.
 - b. Training GANs can be unstable.
 - c. GANs always converge to a globally optimal solution.
 - d. GANs require a large amount of labeled training data.
- ❑ Answer: a, b
- ❑ What is the role of the discriminator in a GAN?
 - a. To generate synthetic data.
 - b. To evaluate the quality of generated samples.
 - c. To provide feedback to the generator.
 - d. To discriminate between real and fake data.
- ❑ Answer: d

